
Fairness Auditing of an Income-Classification Model: SHAP-Based Analysis of Proxy Features

Author One Affiliation One author.one@example.com	Author Two Affiliation Two author.two@example.com	Author Three Affiliation Three author.three@example.com
Author Four Affiliation Four author.four@example.com	Author Five Affiliation Five author.five@example.com	

Abstract

Removing an explicitly protected attribute from a predictive model does not necessarily remove the information associated with that attribute. This report studies that issue in the UCI Adult Income dataset by comparing XGBoost models trained with and without the sensitive attribute `sex`, investigating suspected proxy features with TreeSHAP and feature ablation, and evaluating a Fairlearn ExponentiatedGradient mitigation model. The analysis is intended to separate predictive performance from group-fairness measurements and to make the trade-offs introduced by feature removal and fairness constraints explicit. Empirical conclusions are limited to the dataset, preprocessing choices, model class, and fairness definitions used in this study.

1 Introduction

Machine-learning models are increasingly used to support decisions about people, including decisions related to income, credit, employment, and access to services. In these settings, evaluating predictive performance alone is insufficient: a model can be accurate on average while exhibiting different error rates or prediction rates across demographic groups. Fairness auditing therefore requires both a clear definition of the protected groups and an explicit comparison between predictive and group-level outcomes.

A common first intervention is to remove a sensitive attribute from the model inputs. This intervention is appealing because it is simple and easy to explain, but it does not guarantee that the resulting model is insensitive to that attribute. Other variables may contain information correlated with the protected attribute, either because they reflect social structure or because they encode related aspects of an individual's circumstances. In this report, we refer to such variables as potential *proxy features*. The central question is consequently not only whether `sex` is present in a model, but also whether the remaining feature set allows the model to reconstruct information associated with `sex`.

We investigate this question using the UCI Adult Income dataset, a standard benchmark containing demographic, educational, occupational, and work-related attributes with a binary target indicating whether annual income exceeds \$50,000. The dataset is useful for a controlled audit because it contains an explicit `sex` attribute together with several variables that may encode related social and economic structure. At the same time, it is not a neutral representation of the real world: its categories, missing-value patterns, binary treatment of sex, and historical collection context limit the claims that can be made from any experiment on it. Our conclusions are therefore about this audit design and not about fairness in income prediction in general.

The analysis follows a reproducible comparison of model variants. We first train a baseline XGBoost classifier with `sex` available and compare it with a model that excludes `sex`. We then examine `relationship`, `marital-status`, `occupation`, and `hours-per-week` as suspected proxies. TreeSHAP feature rankings, subgroup comparisons, and feature-ablation experiments are used as diagnostic tools for assessing whether these variables retain predictive information associated with the sensitive attribute. These analyses are exploratory: a high feature importance or a subgroup difference is evidence of association in this dataset, not proof that a feature causes disparate outcomes.

For each model, we report accuracy, balanced accuracy, precision, recall, and F1 together with demographic-parity difference and equalized-odds difference across `sex`. These metrics answer different questions and are not interchangeable. In particular, reducing a statistical-parity difference can come at the cost of predictive performance, while equalized odds focuses on error-rate parity rather than parity of positive predictions. We therefore keep predictive and fairness metrics in separate parts of the comparison rather than treating one aggregate score as a complete definition of fairness. We additionally evaluate Fairlearn’s ExponentiatedGradient procedure under a fairness constraint to examine the trade-off between unconstrained prediction and constrained group outcomes.

This report makes three practical contributions:

- It provides a reproducible audit of whether excluding an explicit sensitive feature changes fairness measurements in an income-classification pipeline.
- It uses model explanations and ablation to test the narrower claim that particular remaining variables may act as proxies, rather than assuming that all correlated variables are interchangeable.
- It presents feature removal and algorithmic mitigation in a common comparison that distinguishes observed performance changes from the limitations of the selected fairness criteria.

The study does not attempt robustness analysis or threshold analysis, and it does not treat statistical parity or equalized odds as sufficient conditions for a fair decision system.

2 Related Work

Fairness definitions and evaluation. Fairness research distinguishes individual-level requirements from group-level statistical criteria. The “fairness through awareness” framework of Dwork et al. [2012] requires similar individuals to receive similar treatment under a task-specific metric, whereas our study evaluates group differences across `sex`. For supervised classification, Hardt et al. [2016] formalized equality of opportunity and equalized odds in terms of conditional prediction rates. These ideas directly motivate our reporting of demographic-parity and equalized-odds differences alongside accuracy and class-balanced performance. Our goal is not to select one metric as a universal definition of fairness, but to show how conclusions about feature removal and mitigation depend on which group-level property is measured.

Proxy features and omitted sensitive attributes. The concern that a model can use protected-group information indirectly is central to proxy-discrimination work. Datta et al. [2017] formalized proxy non-discrimination as a property of data-driven systems in which correlated attributes can influence outputs even when the protected attribute is absent. We operationalize this concern in a narrower empirical setting: we remove `sex`, rank candidate variables with TreeSHAP, compare subgroups, and ablate `relationship`, `marital-status`, `occupation`, and `hours-per-week`. Thus, our contribution is not a new formal definition of proxy discrimination; it is a reproducible audit that connects proxy-focused diagnostics to predictive and fairness outcomes in the same model family. The use of SHAP is diagnostic rather than causal: an important or subgroup-associated feature is treated as a candidate proxy, not as proof that the feature produces discrimination.

Mitigation and subgroup coverage. A second line of work changes the learning procedure instead of relying only on feature deletion. Agarwal et al. [2018] reduced constrained fair classification to cost-sensitive classification problems, providing the algorithmic basis for the ExponentiatedGradient comparison in this report. This lets us place two interventions in one audit: restricting the feature set and imposing a fairness constraint while retaining the broader feature representation. Finally,




Figure 1: EDA of the cleaned Adult data. Left: target distribution. Right: observed proportion above \$50,000 by sex. These are descriptive dataset summaries, not predictions from a trained model.

Buolamwini and Gebru [2018] demonstrated that aggregate performance can conceal substantial intersectional disparities. Our primary audit is across the binary sex labels available in the Adult dataset, so we do not claim to resolve intersectional fairness; instead, we use this prior work to frame subgroup analysis and to make the dataset’s binary attribute representation an explicit limitation.

3 Methodology

3.1 Our proposed method

The manuscript workflow used the local Scientific Agent Skills procedures for evidence-aware scientific writing and consistency checks [Kassis et al., 2026].

3.2 Dataset

We use the UCI Adult Income dataset [Dua and Graff, 2019], comprising 48,842 records and 14 predictors. The binary target is whether annual income exceeds \$50,000, and sex is retained as a sensitive attribute for evaluation. We preserve the dataset’s supplied training and test files rather than drawing a new split.

ETL The ETL implementation reads `adult.data` and `adult.test` from `data/adult`. It removes surrounding whitespace, removes the period appended to labels in the test file, validates the target labels and numeric fields, and converts the target to 0/1. The documented ? values in categorical columns are mapped to an explicit Missing category rather than being silently dropped. This affected 6,465 values: 2,799 in `workclass`, 2,809 in `occupation`, and 857 in `native-country`. The cleaned training and test splits, combined dataset, and ETL metadata are written to `data/processed`.

EDA The exploratory analysis is descriptive and is performed on the cleaned combined dataset before model fitting. It reports descriptive statistics and creates visualizations for the target distribution, age by income class, income rates by education, numeric feature distributions, and documented missing-value counts. We also compare the target rate across the two observed sex categories as an initial descriptive fairness diagnostic; this comparison is not a model fairness result. In the cleaned data, 11,687 of 48,842 records are above \$50,000. The corresponding observed rates are 30.38% for males (9,918/32,650) and 10.93% for females (1,769/16,192).

3.3 Evaluation metrics

We evaluate each model on the held-out test split using both predictive performance and group fairness metrics. Predictive performance is summarized by accuracy, balanced accuracy, precision, recall, and F1 score. Accuracy is the fraction of correct predictions; balanced accuracy is the mean recall across the two income classes; precision is the fraction of predicted above-\$50,000 cases that are correct; recall is the fraction of actual above-\$50,000 cases that are identified; and F1 is the




Figure 2: Observed proportion above \$50,000 by education category in the cleaned dataset.

harmonic mean of precision and recall. Reporting both accuracy and balanced accuracy is useful because the target is not evenly distributed.

Fairness metrics are computed across the observed binary `sex` categories, using the same fixed prediction rule for every model. The demographic-parity difference is defined as

$$\max_{g \in \mathcal{G}} P(\hat{Y} = 1 \mid A = g) - \min_{g \in \mathcal{G}} P(\hat{Y} = 1 \mid A = g),$$

where A is `sex` and $\mathcal{G}$ is the set of observed group labels. It measures disparity in positive prediction rates, without conditioning on the true income label. The equalized-odds difference is reported as the larger of the across-group differences in true-positive rates and false-positive rates:

$$\max \left\{ \max_g \text{TPR}_g - \min_g \text{TPR}_g, \max_g \text{FPR}_g - \min_g \text{FPR}_g \right\}.$$

Thus, it measures disparity in both types of conditional classification outcomes. Lower values indicate smaller observed group differences for both fairness measures, while zero does not establish that a system is fair in all relevant senses.

The comparison table will include the baseline XGBoost model with `sex`, a model with `sex` removed, models with suspected proxy features ablated, and Fairlearn ExponentiatedGradient models constrained by demographic parity or equalized odds. The same predictive and fairness metrics will be reported for each variant so that feature removal, proxy removal, and algorithmic mitigation can be compared without collapsing their different objectives into a single score. We do not perform robustness or threshold analyses; the reported values correspond to the prespecified evaluation procedure and fixed prediction rule.

4 Results

5 Discussion

6 Conclusion

References

- Alekh Agarwal, Alina Beygelzimer, Miroslav Dudík, John Langford, and Hanna Wallach. A reductions approach to fair classification. *arXiv preprint arXiv:1803.02453*, 2018. URL <https://arxiv.org/abs/1803.02453>.
- Joy Buolamwini and Timnit Gebru. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, volume 81

- of *Proceedings of Machine Learning Research*, pages 77–91. PMLR, 2018. URL <https://proceedings.mlr.press/v81/buolamwini18a.html>.
- Anupam Datta, Matt Fredrikson, Gihyuk Ko, Piotr Mardziel, and Shayak Sen. Proxy non-discrimination in data-driven systems. *arXiv preprint arXiv:1707.08120*, 2017. URL <https://arxiv.org/abs/1707.08120>.
- Dheeru Dua and Casey Graff. Uci machine learning repository: Adult dataset, 2019. URL <https://archive.ics.uci.edu/dataset/2/adult>.
- Cynthia Dwork, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. Fairness through awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 214–226, 2012. doi: 10.1145/2090236.2090255. URL <https://arxiv.org/abs/1104.3913>.
- Moritz Hardt, Eric Price, and Nathan Srebro. Equality of opportunity in supervised learning. *arXiv preprint arXiv:1610.02413*, 2016. URL <https://arxiv.org/abs/1610.02413>.
- T. Kassis, V. Agarwal, Y. He, D. Patel, and A. M. Brueckner. Scientific agent skills: A library of procedural knowledge for research agents. *arXiv preprint arXiv:2609.00065*, 2026. doi: 10.48550/arXiv.2609.00065. URL <https://doi.org/10.48550/arXiv.2609.00065>.